

CS562 – The Project

Due Date: 05/08/2014 (Thu.)

Objectives: To become familiar with the concept of *Ad-Hoc OLAP Query Processing*.

Introduction: Your mission in this project is to build a query processing engine for Ad-Hoc OLAP queries. The query construct is based on an extended SQL syntax known as MF and EMF queries (i.e., Multi-Feature and Extended Multi-Feature queries).

PROBLEM: Ad-hoc OLAP queries (also known as multi-dimensional queries) expressed in standard SQL, even the simplest types, often lead to complex relational algebraic expressions with multiple joins, group-bys, and sub-queries. When faced with the challenges of processing such queries, traditional query optimizers do not consider the “big picture”. Rather, they try to optimize a series of joins and group-bys, leading to poor performance.

A SOLUTION: Provide a syntactic framework to allow succinct expression of ad-hoc OLAP queries by extending the *group-by* statement and adding the new clause, *such that*, and in turn, provide a simple, efficient and scalable algorithm to process the queries.

Please refer to the following two research articles for further details on the new syntax and the corresponding processing algorithm:

- “*Querying Multiple Features of Groups in Relational Databases*”, D. Chatziantoniou and K. Ross
- “*Evaluation of Ad Hoc OLAP: In-Place Computation*”, D. Chatziantoniou

The MF and EMF queries take advantage of the grouping variables and such that clause to avoid the need for using multiple sub-queries and joins, and thereby making the expression of the queries more succinct.

Grouping Variables represent subsets of tuples within each group (esp., in MF queries) and the **Such That** clause provides predicates to define the ranges of the grouping variables (i.e., they act as the ‘where’ clauses for the grouping variables). A new relational operator Φ is introduced to capture the essence of the extended SQL (of the grouping variables and the corresponding such that clause). Additionally, there is an evaluation algorithm for the new relational operator that is simple and scalable.

The new relational operator Φ has the following arguments:

- S = List of projected attributes for the query output
- n = Number of grouping variables
- V = List of grouping attributes
- $[F] = \{F_0, F_1, \dots, F_n\}$, list of sets of aggregate functions. F_i represents a list of aggregate functions for each grouping variable.
- $[\sigma] = \{\sigma_0, \sigma_1, \dots, \sigma_n\}$, list of predicates to define the ranges for the grouping variables.
- G = Predicate for the *having* clause (not required for the project)

The evaluation algorithm is as follows, and it utilizes a compact data structure known as '**mf-structure**' – it holds the data corresponding roughly to the output of the MF/EMF query. H represents the mf-structure in the following algorithm. For further details on the 'mf-structure', please refer to the second article, "*Evaluation of Ad Hoc OLAP: In-Place Computation*" (Note that the following algorithm is for the EMF queries, but it also works for the MF queries. The algorithm can be tailored for the MF queries by modifying the third line to "for the entries of H with matching grouping attributes").

```
for scan sc=0 to n {
  for each tuple t on scan {
    for all entries of H,
      check if the defining condition of grouping var
      Xsc is satisfied. If yes, update Xsc's aggregates of the entry
      appropriately.
      X0 denotes the group (the defining condition of X0 is X0.S = S,
      where S denotes the grouping attributes.)
    }
  }
}
```

Project Description:

The following is an outline of the key functionalities of the query processing engine (QPE) you will be implementing for the project:

- The MF queries for the project will be based on the following *schema*
 - sales (cust,prod,day,month,year,state,quant)
The table stores the information about the purchases of a product by a customer on a date and state for a sale amount.
- The *input* for the query processing engine is the list of arguments for the new operator Φ (in place of the actual query represented in SQL) – i.e., you can assume the query has already been transformed into a corresponding relational algebraic expression. For example for the following query,

```
select cust, sum(x.quant), sum(y.quant), sum(z.quant)
from sales
group by cust: x, y, z
such that x.state = 'NY'
        and y.state = 'NJ'
        and z.state = 'CT'
```

you can expect an input such as

```
SELECT ATTRIBUTE(S):
cust, 1_sum_quant, 2_sum_quant, 3_sum_quant
NUMBER OF GROUPING VARIABLES(n):
3
GROUPING ATTRIBUTES(V):
cust
F-VECT([F]):
1_sum_quant, 2_sum_quant, 3_sum_quant
SELECT CONDITION-VECT([σ]):
1.state='NY'
2.state='NJ'
3.state='CT'
```

- Given the input as described above, the query processing engine **generates a program** (written in C, C++, Java, etc.) which implements the evaluation

algorithm mentioned earlier. The generated program (for a given MF query) goes against the `sales` table stored in the PostgreSQL database and generates the output corresponding to the query represented by the input. The generated program can be separately compiled and executed to produce the final output for the input query.

Project Plan: The following is an outline of the project schedule to help you manage your own progress.

Date	Your Query Processing Engine should be able to...
03/26/2014	Read the input and generate a program to create the 'mf-structure' (in memory) – i.e., the generated program contains a list of type/class definitions and variables for the 'mf-structure'.
04/02/2014	<p>Separately generate a program to fulfill a simple query using a cursor. Use the following as a sample query:</p> <pre>select cust, prod, avg(quant), max(quant) from sales where year=2009 group by cust, prod</pre> <p>Make sure the generated program computes the aggregate functions (e.g., avg, max) internally without relying on the PostgreSQL DBMS engine (i.e., the only allowed DB operation is a cursor based scan of the underlying table).</p> <p>The main purpose of this exercise is to help you “visualize” the implementation of the evaluation algorithm in the generated program.</p>
04/16/2014	Generate a program to fulfill a simple MF query that requires a single scan (i.e., an MF query with one grouping variable).
05/02/2014	Generate a program to fulfill MF queries requiring multiple scans.

Grading:

This is a team project – a team consisting of up to 4 members – and the grading of the project will be largely based on the following guideline (similar to other programming assignments you have worked on in the past).

- (80 pts.) Logic/Correctness
- (20 pts.) Programming Style (e.g., comments, indentation, use of functions, etc.). You must include a program header, function header, etc. to clearly state what your program and functions are designed to do. Also for inline comments, please state clearly the purpose of those statements – for you as the programmer and to help others better understand your programming logic.

A program with compilation errors will earn no more than 50 points.

I will schedule a series of sessions for the weeks of 5/3/2014 and 5/10/2014, and your team will be presenting a demo of your projects for 15-20 min. During the demo, you will be asked to input different queries, modify the queries, and in certain cases, modify the code for some simple changes to the algorithms.

***Running the
Query
Processing
Engine***

- Your query processing engine can read the input from a file or interactively from the user.
- The generated program can be manually compiled and executed to generate the final output or the engine itself can directly compile and execute the generated program and output the result of the input query.

Submission:

Please submit your source code (file) with your name and CWID (Campus Wide ID) on it on the course WebCT.

Project Grade Sheet for _____

Major Area	Item	Max	Deduct	Score	%	Total
<i>Compilation</i>	If fails, subtract ...	50				
<i>Logic</i>						
	Total	100			80%	
<i>Style</i>	Header Comment (Overall comments about the project)	10				
	Function Comments	20				
	Line Comment	20				
	Meaningful Names (for functions, variables, etc.)	10				
	Strings – Left Justified	15				
	Numbers – Right Justified	15				
	Modular design (use of classes, methods/functions)	10				
	Total	100			20%	
<i>Total</i>		100			100%	